

Architecture Over Intelligence

Designing reliable AI workflows

Amar Patel · The AI Fellowship Madrid · 9 Oct 2026




It didn't crash. *That was the problem.*



I moved my meeting pipeline off a frontier model onto a cheaper one. It ran all the way through.

It drifted at step 3. By step 5 it had stopped noticing. No exception, no crash, no alert: just plausible output, filed.

A chain hides its own failure



"Investigate this incident" sounds like one job. It is a chain of small jobs, and each one treats the last one's output as fact.

Step 4 trusts step 3. Step 5 trusts step 4. A mistake doesn't stop the chain. It rides it to the end.

My demo never showed me this

In a demo

- One uninterrupted session
- Mistakes are annoying
- A human watches the whole time
- A frontier model

In production

- Compaction, reset, walk-aways
- Mistakes are data loss
- A human checks hours later
- A cheaper model that drifts

A demo works because nothing interrupts it, I'm watching, and I'm running the best model I can buy.

None of those hold on a Tuesday. **Reliability fails at the boundaries**, and no model is big enough to fix a boundary.

Stop letting the model *remember*

FAILURE *It lost its place, and redid work it had already done.*

● process-state.json

meeting_id: 2026-01-15_140000

current_step	2.6
completed_steps	[-1, 0, 1, 2, 2.5, 2.6]
user_confirmed.summary	true
aborted	false

One file on disk holds the position: read before every step, written after.

The model proposes. The file decides.

Stop asking the model for *permission*

FAILURE *Told to present the summary and continue, it continued. Nobody had said yes.*



Auto

Runs without stopping. Nothing here is irreversible.

Reads, health checks, sentiment.



Human

Stops and waits for your yes.

Anything that changes data.



Blocked


Cannot continue. Something required is missing.

Missing file, corrupted state.

A gate isn't asking whether the model got it right. It's asking whether a human said yes. That's authority, not quality control.

Stop trusting the *previous step*

FAILURE *A script failed quietly. The state said it had worked. The next step ran on a file that was never written.*



Every step declares two things: what it produces, and what it needs.

Before it runs, it looks on disk for its inputs. Missing? It stops and says which file. It never guesses.

Assume you will be *interrupted*

FAILURE *The context compacted mid-pipeline. The model forgot everything it was doing.*

step 3 finishes

context reset · you walk away

next run resumes at 4




The file is written after every step, so nothing depends on memory. The reset wiped the model. The file was still on disk.

The next run read it and carried on at step 4. This is the one that made me trust the whole thing.

What it runs, every working day

Transcript in. Speakers named, summary written, actions pulled and triaged, note filed to my vault.



● stops for a human ● runs alone ○ optional

325

MEETINGS PROCESSED

9

HUMAN CHECKPOINTS

8

LLM CALLS

1

STATE FILE


What still broke

- ✎ The state file got edited to the wrong step → *validate on load*
- ⏮ The model skipped a gate → *the file decides, not the prompt*
- ✓ The context reset mid-pipeline → *the file survived, and it resumed where it left off*

Months of daily use, three failures. Each fix moved a decision out of the model and into the file.

Trust the file. Check the artifact itself. Never trust what the last step claims it did.

The model that should have made this *obsolete*



Jev cannot hallucinate, because it never writes text. If anything was going to retire this talk, it was this. Then I read the best write-up of it, titled:

“*The State Machine Is the Agent.*”

Days later researchers bent its decisions with prompt injection. Even a model built to be safe sits behind a deterministic loop, with a human on the irreversible branch.


The same shape fits your work



Every row is the same shape: read, transform, a human approves, write to the system of record. Only the nouns change.

Run these three over your own agent

- 1 What is your single source of truth: a file, a row, a record?
- 2 Where does it stop for a human, and can a step skip that stop?
- 3 Does each step check its inputs exist before it runs?



You don't remove the human as models improve. You move the dial: from checking every step, to guarding the steps that cannot be undone.

Architecture over intelligence

- **The model changes. The structure stays.** A chain breaks wherever nothing supervises it. Put a file in charge.
- **The human doesn't disappear.** Better models move it to the irreversible steps. It cannot delete, overwrite or send without a yes.
- **The same structure runs any model.** When next month's model lands, your workflow doesn't change.

Tonight: add one field, `current_step`. Read it back before every step. Stop if it doesn't match.

Reliability at the boundaries.