


Architecture Over Intelligence

How I made AI workflows reliable, even on a small model

Amar Patel · The AI Fellowship Madrid · 9 October 2026



Make one button blue. Nothing else.



Task: make Add to Cart blue 

Opusfived, a parody by Milos Novovic:
opusfived.dev

It changes both buttons and says "Done".
Every round of clean-up is paid for.



IN SHORT

A clever AI can do more than you asked, call it done, and bill you for the clean-up.

I hit my limit, halfway through a meeting

BEFORE · UNTIL MAY 2026

Claude (cloud)

- Very capable
- Usage limit on my plan
- My meetings leave my house

AFTER · FROM MAY 2026

Qwen 3.6 (open model, on my desk)

- Runs on a 16 GB gaming graphics card
- No limit, no bill
- Private: nothing leaves the house

**“Economics started the journey.
Drift is what I found along the way.**

IN SHORT

I moved to a smaller, local AI to save money. Then it started to go wrong.

Smarter is not the same as correct

drift = the AI slowly goes off track, and does not notice

RUN 1 · QWEN ON MY OWN COMPUTER

- ✗ Wrote the note **before** it knew who was speaking
- ✗ Named **4 of 4** speakers wrong. The CEO became “head of alliances”
- ✗ Deleted the transcript **without asking**

RUN 2 · SAME MODEL, IN THE CLOUD · “SEEMED SMARTER”

- ✗ **Two** notes for one meeting
- ✗ The **wrong date**
- ✗ **2 of 5** quotes given to the wrong person
- ✗ A Chinese word in an English sentence: “metric归属”

Test meeting: “OKR Planning with GitLab Executive Team”, GitLab Unfiltered, public on YouTube.

IN SHORT

Same meeting, my May workflow: both runs went wrong. The smarter one hid it better.

More instructions. Then SHOUTING.

Lines in the prompt (instructions to the AI)



Warnings in CAPITALS (MUST, NEVER, STOP...)



Scripts doing the sure things (normal code)



A REAL LINE FROM 8 SEPTEMBER
If preconditions fail → STOP.
Report gate failure.


It helped a little.

But an instruction is only a suggestion. The model can still decide something else matters more.

What worked: moving the work **out of the prompt** and into code.

IN SHORT If you write MUST in capitals, your process is missing a check.

My whole flow, split three ways



Between every row, a script checks: is it this step's turn, and did the last step leave its file? If not, stop.

IN SHORT The model thinks. Scripts do. Files remember.

Give each step to the right worker

Model

Does this need judgement?

Who is speaking

Summary

Actions

Priorities

Tone of the meeting

Script

Is there only one right answer?

Date and time

File name and folder

Copy the transcript

Link to the project

Delete old files

5,000 words

In May, the model retyped the whole transcript. It ran out of space, and I paid for every word.

Wrong date

The date was already in the file name. The model still got it wrong.


1 → 8 scripts

Today, scripts do the sure things. They get them right every time.

IN SHORT

If the answer never changes, don't ask a model.

Every step leaves a file



A chat forgets

It gets long. Early details get lost. You cannot check it.

My test, 64K context: the harness auto-compacted four times. After each compaction, the model forgot work it had already finished. It did not finish.

A file remembers

The next step starts from the file, not the chat. A script can check it: **is it there? does it have the right shape?**


Same model, same 64K: this time I cleared the context after every step, on purpose. Each step started again from the files. Zero auto-compactions. It finished every step.

IN SHORT

Each step writes down its result, so you can clear the context between steps.

Keep progress outside the model


● progress file .process-state.json	
current_step	2.6
completed_steps	-1, 0, 1, 2, 2.5, 2.6
speakers confirmed	yes
aborted	no



In rehearsal I closed everything in the middle of step 3. When I opened it again, it read this file and carried on. **Nothing repeated. Nothing lost.**

IN SHORT The model forgets everything. The file doesn't.

Check before and after every step



IN SHORT

The check is a small script, not a sentence in the prompt. The model cannot talk it round.

Mark the one-way doors

One-way door · step 7: clean up

Delete the transcript, audio and progress file for this meeting. Approve?

Approve

Block

Only I can answer. The model can't see or click it. Every click is logged.

Retests after Block: 3 of 3 stopped and asked.

AFTER I CLICK BLOCK · PI → MODEL

Amar blocked step 7. Stop and ask what to change. This session can now only read files and run the pipeline scripts.

AND THE SHELL MAKES SURE

-  Delete, move or copy files · edit in place
-  Run code it wrote itself · run any other script
-  Try the same door again
-  Read files · run my pipeline scripts · ask me

IN SHORT

The instruction says stop. The shell makes sure.

My flow in May vs my flow today

SAME MEETING, MODEL AND ANSWERS	MAY		TODAY	
Checked the speakers before writing				
Quotes given to the right person		3 of 5		5 of 5
One note per meeting		2 notes		1 note
Right date				
Asked before deleting				
Carried on after I quit		not tested		

To be fair: today's flow also has more steps and better scripts. This is not a lab test. It is my flow in May against my flow today.

IN SHORT The old flow went wrong quietly. The new one goes wrong loudly, early, and where I can fix it.

Smaller models, bigger jobs

CHEAPER

A smaller model does the judgement. Scripts do the rest.
No tokens spent on work with one right answer.

SAFER

Every step leaves a file. Every Approve and Block is logged.
You can show exactly what happened, and who said yes.

PORTABLE

Swap the model without rewriting the flow.
Less lock-in. Data can stay on your machine, if a model that fits clears the bar.

A SMALLER MODEL. NOT THE SMALLEST · ONE MEETING, FIVE MODELS

One couldn't read the steps. Same Qwen at **3-bit** invented a speaker who was in Vegas. At **8-bit** it got everyone right. **Architecture lowers the bar. It doesn't remove it.**

The same shape fits: sales call follow-ups · invoice processing · customer onboarding · hiring · month-end close

IN SHORT

Reliable flows let you use smaller models on bigger jobs.

Models are splitting: some write, some decide

SPLIT THE MODEL BOX IN TWO

Writer: summary, actions.

Decider: picks one answer from a list.

First out: Jev from TypeSafe AI, which they call a “System One” model. Open models usually catch up within months.

STEPS THAT ONLY CHOOSE

Sentiment

Priority

Project

Follow-up? yes / no

Speaker

The flow offers only legal options. A script checks the answer.

HOW SURE IS IT?



ASK ME

CARRY ON

I am asked only when it is unsure. **The dial belongs to the decider and a script, never to the model doing the work.**

Not wired in yet. This is where it would go.

IN SHORT

Whoever makes the best decider, it slots into one box. The flow stays the same.

Five rules. Start with a checklist.

- 1 Write the steps down,** and give each one to the right worker. Judgement → model. One right answer → script.
- 2 Make every step leave a file.** The next step starts from the file, not the chat.
- 3 Keep progress outside the model.** A note says which step you are on.
- 4 Check before and after every step.** If the answer is no, stop.
- 5 Mark the one-way doors.** A person says yes before anything is sent, deleted or filed.

Try one tonight: take your messiest AI workflow. Write down the steps, and what each one should produce.

IN SHORT Without code, you are the check. Add scripts where the answer never changes, and test them.

Models change every month.

“ Build your flow so the model is a part you can swap.



Thank you

Amar Patel

[linkedin.com/in/amarapatel](https://www.linkedin.com/in/amarapatel)

Scan to connect