

Architecture Over Intelligence


How I made AI workflows reliable, even on a small model

Amar Patel · The AI Fellowship Madrid · 9 October 2026



A GAME

Make one button blue. *Nothing else.*



Opusfived, a parody by Milos Novovic: **opusfived.dev**

The button never turns blue. And every check is paid for. € € €

IN SHORT

A clever AI can waste time and money on a very simple job.

I hit my limit, halfway through a meeting

BEFORE · UNTIL MAY 2026

Claude (cloud)

- Very capable
- Usage limit on my plan
- My meetings leave my house

AFTER · FROM MAY 2026

Qwen 3.6 (open model, on my desk)

- Runs on a 16 GB gaming graphics card
- No limit, no bill
- Private: nothing leaves the house

“Economics started the journey. Drift is what I found along the way.”

IN SHORT

I moved to a smaller, local AI to save money. Then it started to go wrong.

DRIFT

Smarter is not the same as *correct*

drift = the AI slowly goes off track, and does not notice

RUN 1 · QWEN ON MY OWN COMPUTER

- ✗ Wrote the note **before** it knew who was speaking
- ✗ Named **4 of 4** speakers wrong. The CEO became “head of alliances”
- ✗ Deleted the transcript **without asking**

RUN 2 · SAME MODEL, IN THE CLOUD · “SEEMED SMARTER”

- ✗ **Two** notes for one meeting
- ✗ The **wrong date**
- ✗ **2 of 5** quotes given to the wrong person
- ✗ A Chinese word in an English sentence: “metric归属”


Test meeting: “OKR Planning with GitLab Executive Team”, GitLab Unfiltered, public on YouTube.

IN SHORT

Same meeting, my May workflow: both runs went wrong. The smarter one hid it better.

More instructions. Then *SHOUTING*.


Lines in the prompt (instructions to the AI)



Warnings in CAPITALS (MUST, NEVER, STOP...)



Scripts doing the sure things (normal code)



A REAL LINE FROM 8 SEPTEMBER

If preconditions fail → STOP.
Report gate failure.

It helped a little.

But an instruction is only a suggestion. The model can still decide something else matters more.

What worked: moving the work **out of the prompt** and into code.

IN SHORT

If you write MUST in capitals, your process is missing a check.

THE TURN

Split the work into three parts



Model

THINKS · JUDGEMENT

- Who is speaking?
- What was important?
- What did people agree to do?



Scripts

DO · THE SURE THINGS

- Dates and file names
- Filing the note
- Checking every step



Files

REMEMBER · MEMORY

- Which step we are on
- What is already done
- What each step produced

Qwen is good at two things: **calling tools** and **writing clean data (JSON)**. So I gave it only the judgement. I worked this out with help from Claude and Qwen.

IN SHORT

The model thinks. Scripts do. Files remember.

RULE 1

Give each step to *the right worker*

Model

Does this need judgement?

Who is speaking

Summary

Actions

Priorities

Tone of the meeting

Script

Is there only one right answer?

Date and time

File name and folder

Copy the transcript

Link to the project

Delete old files

5,000 words

In May, the model retyped the whole transcript. It ran out of space, and I paid for every word.

Wrong date

The date was already in the file name. The model still got it wrong.

1 → 6 scripts


Today, scripts do the sure things. They get them right every time.

IN SHORT

If the answer never changes, don't ask a model.

RULE 2

Every step leaves *a file*



A chat forgets

It gets long. Early details get lost. You cannot check it.

A file remembers

The next step starts from the file, not the chat. A script can check it: **is it there? does it have the right shape?**


IN SHORT

Each step writes down its result, so the next step never has to remember.

RULE 3

Keep progress *outside* the model

● progress file	.process-state.json
current_step	2.6
completed_steps	-1, 0, 1, 2, 2.5, 2.6
speakers confirmed	yes
aborted	no



In rehearsal I closed everything in the middle of step 3. When I opened it again, it read this file and carried on. **Nothing repeated. Nothing lost.**


The same file let me run the demos on my local Qwen and on a cloud copy **without changing a line.**

IN SHORT

The model forgets everything. The file doesn't.

RULE 4

Check *before* and *after* every step



IN SHORT

The check is a small script, not a sentence in the prompt. The model cannot talk it round.

RULE 5

Mark the *one-way doors*

One-way door · step 7: clean up

Delete the transcript, audio and progress file for this meeting. Approve?

Approve

Block

Only I can answer. The model can't see or click it. Every click is logged.

Retests after Block: 3 of 3 stopped and asked.

AFTER I CLICK BLOCK · PI → MODEL

Amar blocked step 7. Stop and ask what to change. This session can now only read files and run the pipeline scripts.

AND THE SHELL MAKES SURE

- ✗ Delete, move or copy files · edit in place
- ✗ Run code it wrote itself · run any other script
- ✗ Try the same door again
- ✓ Read files · run my pipeline scripts · ask me

IN SHORT

The instruction says stop. The shell makes sure.

THE RESULT

My flow in May vs *my flow today*

SAME MEETING, MODEL AND ANSWERS	MAY	TODAY
Checked the speakers before writing	✗	✓
Quotes given to the right person	✗ 3 of 5	✓ 5 of 5
One note per meeting	✗ 2 notes	✓ 1 note
Right date	✗	✓
Asked before filing or deleting	✗	✓
Carried on after I quit	– not tested	✓

To be fair: today's flow also has more steps and better scripts. This is not a lab test. It is my flow in May against my flow today.

IN SHORT

The old flow went wrong quietly. The new one goes wrong loudly, early, and where I can fix it.

WHAT THIS UNLOCKS

Smaller models, *bigger jobs*

CHEAPER

A small, local model does the judgement. Scripts do the rest.

No tokens spent on work with one right answer.

SAFER

Every step leaves a file. Every Approve and Block is logged.

You can show exactly what happened, and who said yes.

PORTABLE

Swap the model without rewriting the flow.

Less lock-in. Private data can stay on your own machine.

The same steps, files and doors fit:

Sales call follow-ups

Invoice processing

Customer onboarding

Hiring

Month-end close

IN SHORT

Reliable flows let you use smaller models on bigger jobs.

A new kind of model arrived *this month*

SPLIT THE MODEL BOX IN TWO

Writer: summary, actions.

Decider: picks one answer from a list.

Jev (TypeSafe AI) only decides: a choice, and how sure it is.

STEPS THAT ONLY CHOOSE

Sentiment

Priority

Project

Follow-up? yes / no

Duplicate? yes / no

Speaker

The flow offers only legal options.
A script checks the answer.

HOW SURE IS IT?



ASK ME

CARRY ON

I am asked only when it is unsure.
Not wired in yet. This is where it would go.

IN SHORT

A new model slots into one box. The flow stays the same.

TAKE THESE HOME

Five rules. *Start with a checklist.*

- 1 Write the steps down,** and give each one to the right worker. Judgement → model. One right answer → script.
- 2 Make every step leave a file.** The next step starts from the file, not the chat.
- 3 Keep progress outside the model.** A note says which step you are on.
- 4 Check before and after every step.** If the answer is no, stop.
- 5 Mark the one-way doors.** A person says yes before anything is sent, deleted or filed.

Try one tonight: take your messiest AI workflow. Write down the steps, and what each one should produce.

IN SHORT

Without code, you are the check. Add scripts where the answer never changes, and test them.

Models change every month.

“Build your flow so the model is a part you can swap.”



Thank you

Amar Patel

[linkedin.com/in/amarapatel](https://www.linkedin.com/in/amarapatel)

Scan to connect